

prabhudayala@gmail.com_13

July 22, 2019

```
[1]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.  
      →py  
  
from __future__ import print_function  
import keras  
from keras.datasets import mnist  
from keras.models import Sequential  
from keras.layers import Dense, Dropout, Flatten  
from keras.layers import Conv2D, MaxPooling2D  
from keras import backend as K  
from keras.layers.normalization import BatchNormalization  
  
batch_size = 512  
num_classes = 10  
epochs = 5  
  
# input image dimensions  
img_rows, img_cols = 28, 28  
  
# the data, split between train and test sets  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
  
if K.image_data_format() == 'channels_first':  
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)  
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)  
    input_shape = (1, img_rows, img_cols)  
else:  
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)  
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)  
    input_shape = (img_rows, img_cols, 1)  
  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
x_train /= 255  
x_test /= 255
```

```

print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

Using TensorFlow backend.

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

[2]: import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()

```

```

[3]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,

```

```
batch_size=batch_size,  
epochs=epochs,  
verbose=1,  
validation_data=(x_test, y_test))
```

WARNING: Logging before flag parsing goes to stderr.

W0717 21:39:39.668090 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:74: The name tf.get_default_graph
is deprecated. Please use tf.compat.v1.get_default_graph instead.

W0717 21:39:39.680058 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:517: The name tf.placeholder is
deprecated. Please use tf.compat.v1.placeholder instead.

W0717 21:39:39.682081 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:4138: The name tf.random_uniform is
deprecated. Please use tf.random.uniform instead.

W0717 21:39:39.701999 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:3976: The name tf.nn.max_pool is
deprecated. Please use tf.nn.max_pool2d instead.

W0717 21:39:39.703994 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:133: The name
tf.placeholder_with_default is deprecated. Please use
tf.compat.v1.placeholder_with_default instead.

W0717 21:39:39.711008 8044 deprecation.py:506] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from
tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed
in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 -
keep_prob`.

W0717 21:39:39.811706 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated.
Please use tf.compat.v1.train.Optimizer instead.

W0717 21:39:39.816720 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-

packages\keras\backend\tensorflow_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.

W0717 21:39:39.882516 8044 deprecation.py:323] From C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version. Instructions for updating: Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 64s 1ms/step - loss: 0.5720 - acc: 0.8175 - val_loss: 0.1035 - val_acc: 0.9649

Epoch 2/5

60000/60000 [=====] - 67s 1ms/step - loss: 0.1262 - acc: 0.9628 - val_loss: 0.0631 - val_acc: 0.9772

Epoch 3/5

60000/60000 [=====] - 70s 1ms/step - loss: 0.0880 - acc: 0.9740 - val_loss: 0.0415 - val_acc: 0.9858

Epoch 4/5

60000/60000 [=====] - 69s 1ms/step - loss: 0.0691 - acc: 0.9799 - val_loss: 0.0384 - val_acc: 0.9870

Epoch 5/5

60000/60000 [=====] - 69s 1ms/step - loss: 0.0569 - acc: 0.9835 - val_loss: 0.0472 - val_acc: 0.9836

```
[4]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

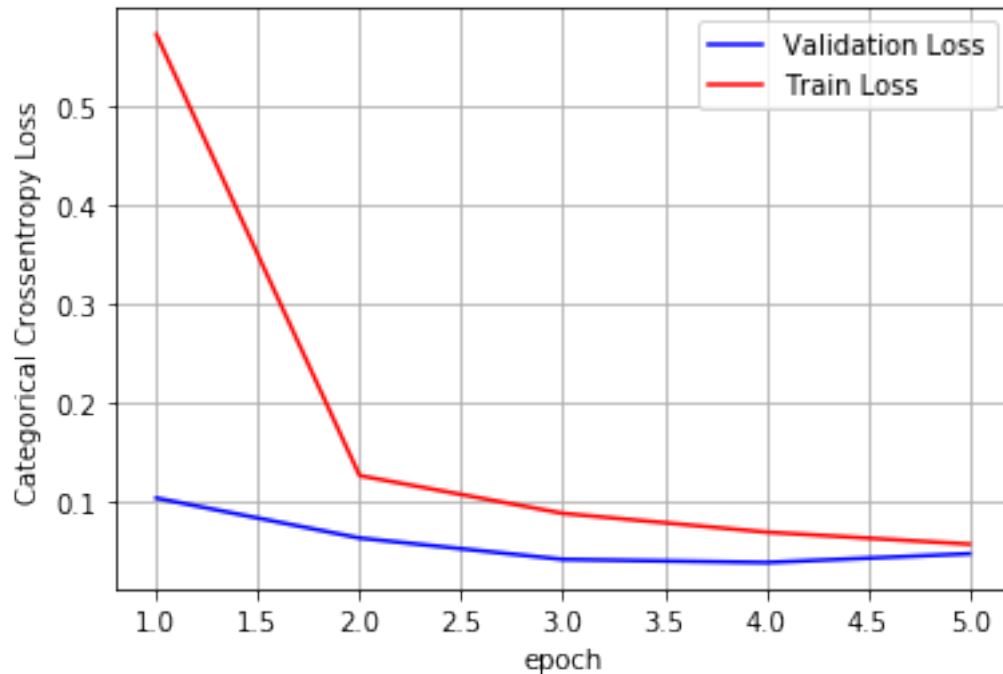
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.04723946803898434

Test accuracy: 0.9836



```
[5]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 58s 960us/step - loss: 0.5140 -

```

acc: 0.8404 - val_loss: 0.1440 - val_acc: 0.9570
Epoch 2/5
60000/60000 [=====] - 61s 1ms/step - loss: 0.1436 -
acc: 0.9574 - val_loss: 0.0626 - val_acc: 0.9802
Epoch 3/5
60000/60000 [=====] - 58s 965us/step - loss: 0.1025 -
acc: 0.9697 - val_loss: 0.0500 - val_acc: 0.9838
Epoch 4/5
60000/60000 [=====] - 57s 957us/step - loss: 0.0816 -
acc: 0.9759 - val_loss: 0.0534 - val_acc: 0.9816
Epoch 5/5
60000/60000 [=====] - 54s 905us/step - loss: 0.0705 -
acc: 0.9788 - val_loss: 0.0367 - val_acc: 0.9877

```

```

[6]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

x = list(range(1,epochs+1))

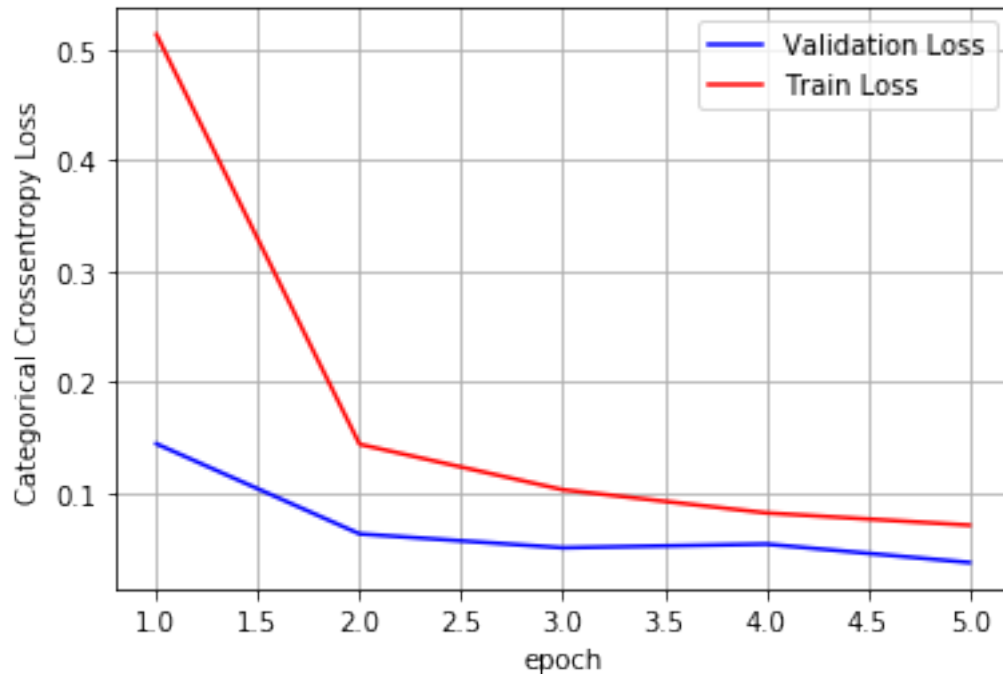
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.03671102816282073
Test accuracy: 0.9877

```



```
[7]: model = Sequential()
model.add(Conv2D(128, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(BatchNormalization())
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
```

```
validation_data=(x_test, y_test))
```

```
W0717 21:50:12.915995 8044 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:1834: The name
tf.nn.fused_batch_norm is deprecated. Please use
tf.compat.v1.nn.fused_batch_norm instead.
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 245s 4ms/step - loss: 0.2719 -
acc: 0.9372 - val_loss: 0.0902 - val_acc: 0.9773

Epoch 2/5

60000/60000 [=====] - 244s 4ms/step - loss: 0.0584 -
acc: 0.9872 - val_loss: 0.0449 - val_acc: 0.9893

Epoch 3/5

60000/60000 [=====] - 244s 4ms/step - loss: 0.0352 -
acc: 0.9919 - val_loss: 0.0397 - val_acc: 0.9886

Epoch 4/5

60000/60000 [=====] - 540s 9ms/step - loss: 0.0249 -
acc: 0.9939 - val_loss: 0.0343 - val_acc: 0.9903

Epoch 5/5

60000/60000 [=====] - 618s 10ms/step - loss: 0.0184 -
acc: 0.9954 - val_loss: 0.0327 - val_acc: 0.9895

```
[8]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

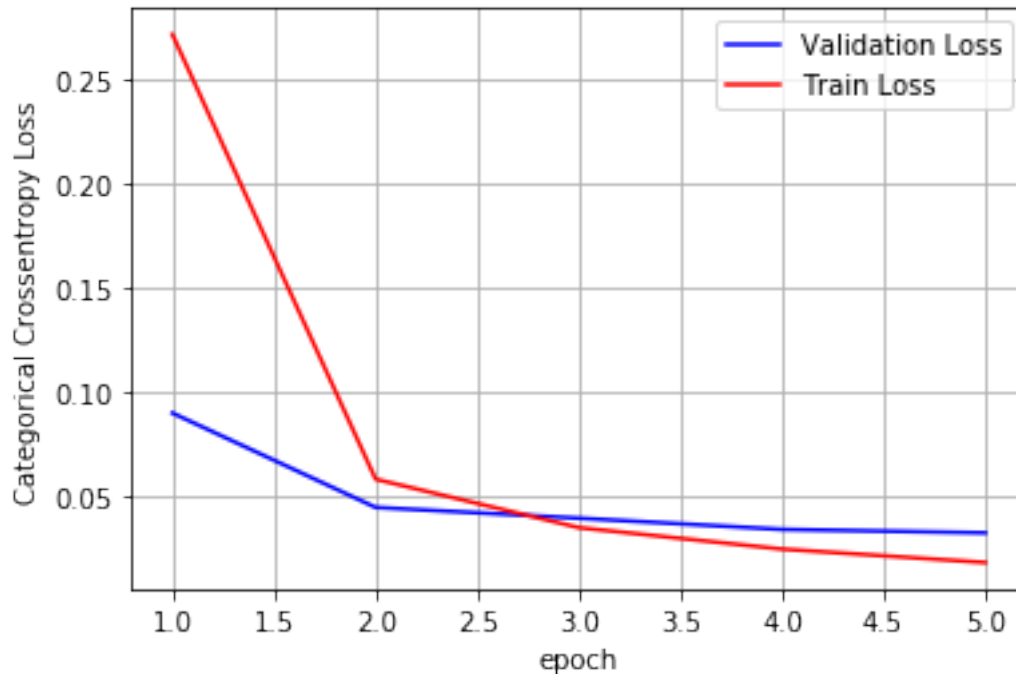
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.03268356256503612

Test accuracy: 0.9895



This model converged very fast than other model.

```
[9]: model = Sequential()
model.add(Conv2D(64, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (5, 5), activation='relu'))
model.add(Dropout(0.25))
model.add(Conv2D(16, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Dropout(0.4))
model.add(BatchNormalization())

model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
```

```
batch_size=batch_size,  
epochs=epochs,  
verbose=1,  
validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 84s 1ms/step - loss: 0.6515 -
acc: 0.8001 - val_loss: 0.1776 - val_acc: 0.9513

Epoch 2/5

60000/60000 [=====] - 83s 1ms/step - loss: 0.1761 -
acc: 0.9476 - val_loss: 0.0963 - val_acc: 0.9704

Epoch 3/5

60000/60000 [=====] - 81s 1ms/step - loss: 0.1195 -
acc: 0.9637 - val_loss: 0.1092 - val_acc: 0.9649

Epoch 4/5

60000/60000 [=====] - 81s 1ms/step - loss: 0.0958 -
acc: 0.9713 - val_loss: 0.0569 - val_acc: 0.9803

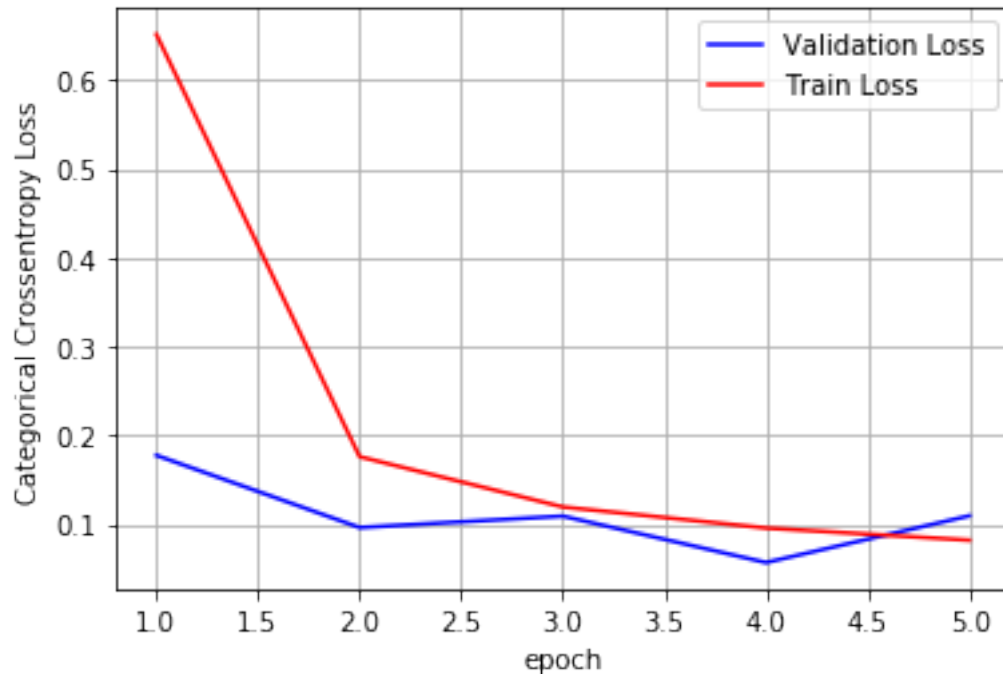
Epoch 5/5

60000/60000 [=====] - 80s 1ms/step - loss: 0.0821 -
acc: 0.9740 - val_loss: 0.1095 - val_acc: 0.9625

```
[10]: score = model.evaluate(x_test, y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
x = list(range(1,epochs+1))  
  
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1095163794439286

Test accuracy: 0.9625



The abnormality in test score might be caused by dropout just after two layers of convolution. May be using Using 3x3 kernel after 5x5 kernel hurts the model.

```
[11]: model = Sequential()
model.add(Conv2D(64, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Dropout(0.4))
model.add(BatchNormalization())

model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                  batch_size=batch_size,
```

```
epochs=epochs,  
verbose=1,  
validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 78s 1ms/step - loss: 0.3386 -
acc: 0.8995 - val_loss: 0.1138 - val_acc: 0.9630

Epoch 2/5

60000/60000 [=====] - 73s 1ms/step - loss: 0.0945 -
acc: 0.9720 - val_loss: 0.0652 - val_acc: 0.9798

Epoch 3/5

60000/60000 [=====] - 75s 1ms/step - loss: 0.0698 -
acc: 0.9789 - val_loss: 0.0426 - val_acc: 0.9875

Epoch 4/5

60000/60000 [=====] - 71s 1ms/step - loss: 0.0559 -
acc: 0.9829 - val_loss: 0.0366 - val_acc: 0.9878

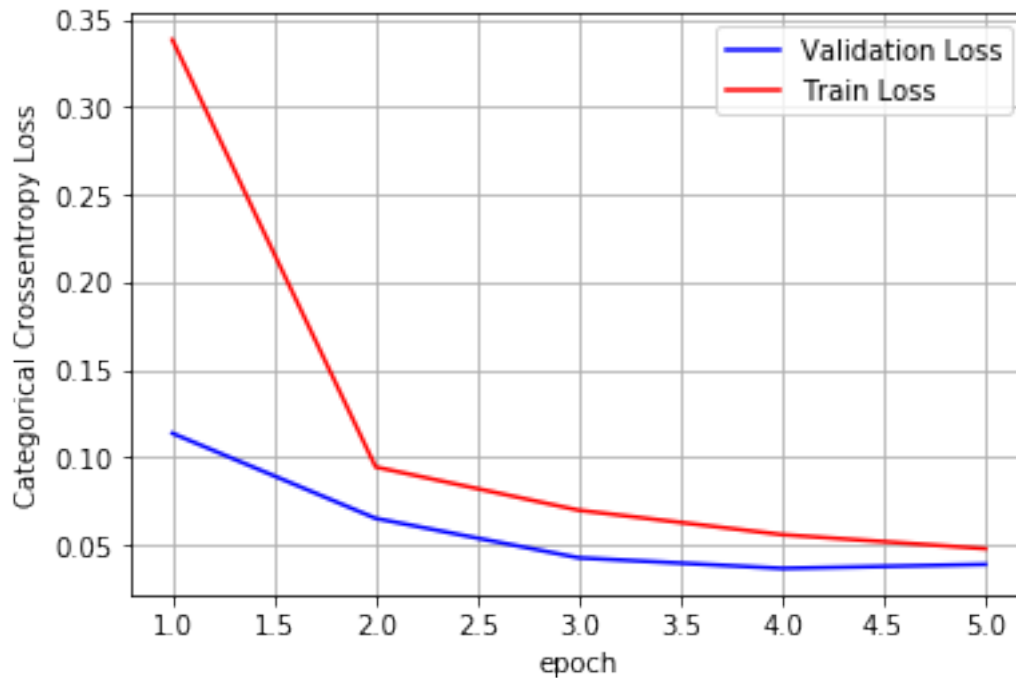
Epoch 5/5

60000/60000 [=====] - 75s 1ms/step - loss: 0.0479 -
acc: 0.9851 - val_loss: 0.0389 - val_acc: 0.9880

```
[12]: score = model.evaluate(x_test, y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
x = list(range(1,epochs+1))  
  
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.03889234391124919

Test accuracy: 0.988



After removing the dropout from the first layer the model performs better.

[13]: `from prettytable import PrettyTable`

```
x = PrettyTable()
x.field_names = ["Model", "Convolution", "MaxPooling", "Dropout", "Batch_
    →Normalization", "optimizer", "Train Accuracy", "Test Accuracy"]
x.add_row(["Model 1", "(2x2)", "(2x2)", "Yes", "Yes", "Adadelata",0.983,0.983])
x.add_row(["Model 2", "(3x3)", "(2x2)", "Yes", "No", "Adadelata",0.978,0.987])
x.add_row(["Model 3", "(3x3) (5x5) (3x3)", "(2x2)", "Yes", "Yes", "Adadelata",0.
    →995,0.989])
x.add_row(["Model 4", "(3x3) (5x5) (3x3)", "(2x2)", "Yes", "Yes", "Adadelata",0.
    →974,0.962])
x.add_row(["Model 5", "(3x3) (5x5)", "(2x2)", "Yes", "Yes", "Adadelata",0.985,0.
    →988])
x.border=True
print(x)
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Model | Convolution | MaxPooling | Dropout | Batch Normalization |
optimizer | Train Accuracy | Test Accuracy |
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Model 1 | (2x2) | (2x2) | Yes | Yes |
Adadelata | 0.983 | 0.983 |
```

Model 2	(3x3)		(2x2)		Yes		No	
Adadelta	0.978		0.987					
Model 3	(3x3) (5x5) (3x3)		(2x2)		Yes		Yes	
Adadelta	0.995		0.989					
Model 4	(3x3) (5x5) (3x3)		(2x2)		Yes		Yes	
Adadelta	0.974		0.962					
Model 5	(3x3) (5x5)		(2x2)		Yes		Yes	
Adadelta	0.985		0.988					
+-----+-----+-----+-----+-----+-----+-----+								
-----+-----+-----+-----+								