

prabhudayala@gmail.com_FB_Models

November 24, 2019

Social network Graph Link Prediction - Facebook Challenge

```
[1]: #Importing Libraries
# please do go through this python notebook:
import warnings
warnings.filterwarnings("ignore")

import csv
import pandas as pd#pandas to create small dataframes
import datetime #Convert to unix time
import time #Convert to unix time
# if numpy is not installed already : pip3 install numpy
import numpy as np#Do arithmetic operations on arrays
# matplotlib: used to plot graphs
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns#Plots
from matplotlib import rcParams#Size of plots
from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
import math
import pickle
import os
# to install xgboost: pip3 install xgboost
import xgboost as xgb

import warnings
import networkx as nx
import pdb
import pickle
from pandas import HDFStore, DataFrame
from pandas import read_hdf
from scipy.sparse.linalg import svds, eigs
import gc
from tqdm import tqdm
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
```

```
[2]: #reading
from pandas import read_hdf
df_final_train = read_hdf('data/fea_sample/storage_sample_stage6.h5',
    ↳ 'train_df',mode='r')
df_final_test = read_hdf('data/fea_sample/storage_sample_stage6.h5',
    ↳ 'test_df',mode='r')

[3]: df_final_train.columns

[3]: Index(['source_node', 'destination_node', 'indicator_link',
    'jaccard_followers', 'jaccard_followees', 'cosine_followers',
    'cosine_followees', 'num_followers_s', 'num_followees_s',
    'num_followees_d', 'inter_followers', 'inter_followees', 'adar_index',
    'follows_back', 'same_comp', 'shortest_path', 'weight_in', 'weight_out',
    'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4', 'page_rank_s',
    'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d', 'authorities_s',
    'authorities_d', 'svd_dot_u', 'svd_dot_v', 'preferential_attachment_p',
    'preferential_attachment_s'],
    dtype='object')

[4]: y_train = df_final_train.indicator_link
y_test = df_final_test.indicator_link


[5]: df_final_train.drop(['source_node',
    ↳ 'destination_node', 'indicator_link'],axis=1,inplace=True)
df_final_test.drop(['source_node',
    ↳ 'destination_node', 'indicator_link'],axis=1,inplace=True)

[6]: estimators = [10,50,100,250,450]
train_scores = []
test_scores = []
for i in estimators:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None,
    ↳ criterion='gini',
        max_depth=5, max_features='auto', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=52, min_samples_split=120,
        min_weight_fraction_leaf=0.0, n_estimators=i,
    ↳ n_jobs=-1,random_state=25,verbose=0,warm_start=False)
    clf.fit(df_final_train,y_train)
    train_sc = f1_score(y_train,clf.predict(df_final_train))
    test_sc = f1_score(y_test,clf.predict(df_final_test))
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Estimators = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(estimators,train_scores,label='Train Score')
plt.plot(estimators,test_scores,label='Test Score')
plt.xlabel('Estimators')
plt.ylabel('Score')
```

```
plt.title('Estimators vs score at depth of 5')
```

```
Estimators = 10 Train Score 0.9253332491662195 test Score 0.9235924932975871
Estimators = 50 Train Score 0.9258777360391177 test Score 0.9238324060023605
Estimators = 100 Train Score 0.9254051778572931 test Score 0.9224591584419976
Estimators = 250 Train Score 0.9248103689520584 test Score 0.9217240577504388
Estimators = 450 Train Score 0.9241257602655546 test Score 0.9207803724504877
```

[6]: Text(0.5, 1.0, 'Estimators vs score at depth of 5')



```
[7]: depths = [3,9,11,15,20,35,50,70,130]
train_scores = []
test_scores = []
for i in depths:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None,
    criterion='gini',
    max_depth=i, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=52, min_samples_split=120,
    min_weight_fraction_leaf=0.0, n_estimators=115,
    n_jobs=-1, random_state=25, verbose=0, warm_start=False)
    clf.fit(df_final_train, y_train)
    train_sc = f1_score(y_train, clf.predict(df_final_train))
```

```


test_sc = f1_score(y_test,clf.predict(df_final_test))
test_scores.append(test_sc)
train_scores.append(train_sc)
print('depth = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(depths,train_scores,label='Train Score')
plt.plot(depths,test_scores,label='Test Score')
plt.xlabel('Depth')
plt.ylabel('Score')
plt.title('Depth vs score at depth of 5 at estimators = 115')
plt.show()

```

```

depth = 3 Train Score 0.917830714226324 test Score 0.9173315510255232
depth = 9 Train Score 0.960628960242471 test Score 0.9248893572181244
depth = 11 Train Score 0.963722253263575 test Score 0.925804278255011
depth = 15 Train Score 0.9665299159185995 test Score 0.9268498407945511
depth = 20 Train Score 0.9666257145468022 test Score 0.9265477218629538
depth = 35 Train Score 0.9666737724721097 test Score 0.9268981208393022
depth = 50 Train Score 0.9666737724721097 test Score 0.9268981208393022
depth = 70 Train Score 0.9666737724721097 test Score 0.9268981208393022
depth = 130 Train Score 0.9666737724721097 test Score 0.9268981208393022

```



```

[8]: from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score

```

```

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

param_dist = {"n_estimators": sp_randint(105,125),
              "max_depth": sp_randint(10,15),
              "min_samples_split": sp_randint(110,190),
              "min_samples_leaf": sp_randint(25,65)}

clf = RandomForestClassifier(random_state=25,n_jobs=-1)

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               ↵
                               ↪n_iter=5,cv=10,scoring='f1',random_state=25,return_train_score=True)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])

```

mean test scores [0.96454843 0.96396638 0.96254065 0.96390786 0.96553771]
mean train scores [0.96591799 0.96493352 0.96334515 0.96515447 0.96724712]

[9]: `print(rf_random.best_estimator_)`

```

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                       max_depth=14, max_features='auto', max_leaf_nodes=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=28, min_samples_split=111,
                       min_weight_fraction_leaf=0.0, n_estimators=121,
                       n_jobs=-1, oob_score=False, random_state=25, verbose=0,
                       warm_start=False)

```

[10]: `clf = RandomForestClassifier(bootstrap=True, class_weight=None, ↵
↪criterion='gini',
max_depth=14, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=28, min_samples_split=111,
min_weight_fraction_leaf=0.0, n_estimators=121, n_jobs=-1,
oob_score=False, random_state=25, verbose=0, warm_start=False)`

[11]: `clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)`

[12]: `from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))`

Train f1 score 0.9679107505070993
Test f1 score 0.9273340358271865

```
[13]: from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1))).T)

    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
    →yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")


    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
    →yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
    →yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")


    plt.show()
```

```
[14]: print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```


Train confusion_matrix




Test confusion_matrix



```
[15]: from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```



```
[16]: features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

1 Assignments:

1. Add another feature called Preferential Attachment with followers and followees data of vertex. you can check about Preferential Attachment in below link <http://be.amazd.com/link-prediction/>
2. Add feature called svd_dot. you can calculate svd_dot as Dot product between source node svd and destination node svd features. you can read about this in below pdf https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf
3. Tune hyperparameters for XG boost with all these features and check the error metric.

```
[17]: from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

param_dist = {
    'max_depth':range(3,10,2),
    'min_child_weight':range(1,6,2),
    'gamma':[i/10.0 for i in range(0,5)],
    'subsample':[i/10.0 for i in range(6,10)],
    'colsample_bytree':[i/10.0 for i in range(6,10)],
    'reg_alpha':[1e-5, 1e-2, 0.1, 1, 100]
}

clf = xgb.XGBClassifier()

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               ↵
                               n_iter=10,cv=10,scoring='f1',random_state=25,return_train_score=True,verbose=1,n_jobs=4)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed: 1.5min
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed: 4.6min finished
```

```
mean test scores [0.9732647  0.97300458 0.97330321 0.97281577 0.97739804
0.97315316
0.97741073 0.97336856 0.97759844 0.9777435 ]
mean train scores [0.97368159 0.97329602 0.97354947 0.97303076 0.98352133
0.97335821
0.98397578 0.97349323 0.98407541 0.98462509]
```

```
[18]: print(rf_random.best_estimator_)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.8, gamma=0.3,
              learning_rate=0.1, max_delta_step=0, max_depth=9,
              min_child_weight=3, missing=None, n_estimators=100, n_jobs=1,
              nthread=None, objective='binary:logistic', random_state=0,
```

```
reg_alpha=0.01, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=0.7, verbosity=1)
```

```
[21]: clf = xgb.XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=0.8, gamma=0.3,
    learning_rate=0.1, max_delta_step=0, max_depth=9,
    min_child_weight=3, missing=None, n_estimators=100, n_jobs=1,
    nthread=None, objective='binary:logistic', random_state=0,
    reg_alpha=0.01, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=0.7)
```

```
[22]: clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)
```

```
[23]: from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))
```

Train f1 score 0.9842347560668486

Test f1 score 0.9270541672806888

```
[24]: from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1)))) .T)

    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
    yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")


    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
    yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")
```

```
plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
→ yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")


plt.show()
```

```
[25]: print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```

Train confusion_matrix




Test confusion_matrix




```
[26]: from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
```

```
plt.legend()
plt.show()
```



```
[27]: features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



```
[29]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "f1score Train", "f1score Test"]
x.add_row(["Random forest", round(0.9679107505070993,3), round(0.
→9273340358271865,3)])
x.add_row(["XGBoost", round(0.9842347560668486,3), round(0.
→9270541672806888,3)])
x.border=True
print(x)
```

```
+-----+-----+-----+
|      Model      | f1score Train | f1score Test |
+-----+-----+-----+
```

Random forest	0.968		0.927	
XGBoost	0.984		0.927	
+-----+-----+-----+				

Due to cold start problem the test accuracy is not up to the mark as train. But certainly XGBoost is the best model for this scenario.