

prabhudayala@gmail.com_12

July 15, 2019

0.1 Keras – MLPs on MNIST

```
[1]: # if you keras is not using tensorflow as backend set
      ↳ "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
[2]: import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
[3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
[4]: print("Number of training examples :", X_train.shape[0], "and each image is of
      ↳ shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of
      ↳ shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
[5]: # if you observe the input shape its 2 dimensional vector
      # for each image we have a (28*28) vector
```

```
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
[6]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (", X_train.shape[1], ")")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (", X_test.shape[1], ")")
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
[7]: # An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253]
```

253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	24	114	221	253	253	253
253	201	78	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	23	66	213	253	253	253	253	198	81	2	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	18	171	219	253	253	253	253	195
80	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
55	172	226	253	253	253	253	244	133	11	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	136	253	253	253	212	135	132	16
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0]							

```
[8]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the
# data
# 
$$X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$$


X_train = X_train/255
X_test = X_test/255
```

```
[9]: # example data point after normalizing
print(X_train[0])
```

[illegible]

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

```
[10]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# ↪ 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```
[11]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
[12]: from keras.models import Sequential
      from keras.layers import Dense, Activation
      from keras.layers import Dropout
      from keras.layers.normalization import BatchNormalization
```

Softmax classifier

Model 1. 2 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[13]: `# https://stackoverflow.com/questions/34716454/
↪ where-do-i-call-the-batchnormalization-function-in-keras`

```

model_1 = Sequential()

model_1.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None),
    ↳input_shape=(input_dim,)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.5))

model_1.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.5))

model_1.add(Dense(output_dim, activation='softmax'))

model_1.summary()

```

WARNING: Logging before flag parsing goes to stderr.
W0715 18:42:29.220162 11376 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:74: The name tf.get_default_graph
is deprecated. Please use tf.compat.v1.get_default_graph instead.

W0715 18:42:29.232118 11376 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:517: The name tf.placeholder is
deprecated. Please use tf.compat.v1.placeholder instead.

W0715 18:42:29.234127 11376 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:4115: The name tf.random_normal is
deprecated. Please use tf.random.normal instead.

W0715 18:42:29.294940 11376 deprecation_wrapper.py:119] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:133: The name
tf.placeholder_with_default is deprecated. Please use
tf.compat.v1.placeholder_with_default instead.

W0715 18:42:29.310896 11376 deprecation.py:506] From
C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-
packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from
tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed
in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

W0715 18:42:29.391703 11376 deprecation_wrapper.py:119] From C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-packages\keras\backend\tensorflow_backend.py:4138: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 128)	100480
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 64)	8256
batch_normalization_2 (Batch Normalization)	(None, 64)	256
dropout_2 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 10)	650
Total params: 110,154		
Trainable params: 109,770		
Non-trainable params: 384		

```
[14]: model_1.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    ↪ verbose=1, validation_data=(X_test, Y_test))
```

W0715 18:42:31.010054 11376 deprecation_wrapper.py:119] From C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

W0715 18:42:31.090867 11376 deprecation.py:323] From C:\Users\user\Anaconda3\envs\tensorflow_cpu\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use `tf.where` in 2.0, which has the same broadcast rule as `np.where`

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 30us/step - loss: 0.9239 -
acc: 0.7141 - val_loss: 0.2713 - val_acc: 0.9185

Epoch 2/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4345 -
acc: 0.8711 - val_loss: 0.2065 - val_acc: 0.9355

Epoch 3/20

60000/60000 [=====] - 1s 22us/step - loss: 0.3440 -
acc: 0.8991 - val_loss: 0.1721 - val_acc: 0.9464

Epoch 4/20

60000/60000 [=====] - 1s 22us/step - loss: 0.2976 -
acc: 0.9140 - val_loss: 0.1457 - val_acc: 0.9540

Epoch 5/20

60000/60000 [=====] - 1s 23us/step - loss: 0.2635 -
acc: 0.9225 - val_loss: 0.1367 - val_acc: 0.9558

Epoch 6/20

60000/60000 [=====] - 1s 22us/step - loss: 0.2377 -
acc: 0.9308 - val_loss: 0.1243 - val_acc: 0.9610

Epoch 7/20

60000/60000 [=====] - 1s 21us/step - loss: 0.2229 -
acc: 0.9361 - val_loss: 0.1137 - val_acc: 0.9641

Epoch 8/20

60000/60000 [=====] - 1s 21us/step - loss: 0.2063 -
acc: 0.9394 - val_loss: 0.1085 - val_acc: 0.9660

Epoch 9/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1965 -
acc: 0.9433 - val_loss: 0.1033 - val_acc: 0.9676

Epoch 10/20

60000/60000 [=====] - 1s 22us/step - loss: 0.1855 -
acc: 0.9469 - val_loss: 0.0983 - val_acc: 0.9708

Epoch 11/20

60000/60000 [=====] - 1s 21us/step - loss: 0.1790 -
acc: 0.9478 - val_loss: 0.0957 - val_acc: 0.9707

Epoch 12/20

60000/60000 [=====] - 1s 21us/step - loss: 0.1723 -
acc: 0.9490 - val_loss: 0.0950 - val_acc: 0.9719

Epoch 13/20

60000/60000 [=====] - 1s 21us/step - loss: 0.1640 -
acc: 0.9512 - val_loss: 0.0941 - val_acc: 0.9716

Epoch 14/20

60000/60000 [=====] - 1s 21us/step - loss: 0.1575 -
acc: 0.9540 - val_loss: 0.0905 - val_acc: 0.9726

Epoch 15/20

60000/60000 [=====] - 1s 21us/step - loss: 0.1534 -

```

acc: 0.9549 - val_loss: 0.0879 - val_acc: 0.9734
Epoch 16/20
60000/60000 [=====] - 1s 21us/step - loss: 0.1511 -
acc: 0.9555 - val_loss: 0.0853 - val_acc: 0.9741
Epoch 17/20
60000/60000 [=====] - 1s 21us/step - loss: 0.1459 -
acc: 0.9570 - val_loss: 0.0907 - val_acc: 0.9736
Epoch 18/20
60000/60000 [=====] - 2s 25us/step - loss: 0.1450 -
acc: 0.9575 - val_loss: 0.0861 - val_acc: 0.9740
Epoch 19/20
60000/60000 [=====] - 1s 21us/step - loss: 0.1356 -
acc: 0.9597 - val_loss: 0.0833 - val_acc: 0.9756
Epoch 20/20
60000/60000 [=====] - 1s 21us/step - loss: 0.1334 -
acc: 0.9597 - val_loss: 0.0806 - val_acc: 0.9764

```

```

[15]: score = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
→number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.08059193601340521
Test accuracy: 0.9764

```



Model 1. 2 hidden layer, ReLU, Dropout and AdamOptimizer

[16]: # [https://stackoverflow.com/questions/34716454/](https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras)
 →where-do-i-call-the-batchnormalization-function-in-keras

```
model_1 = Sequential()

model_1.add(Dense(128, activation='relu',
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None),
    →input_shape=(input_dim,)) )
model_1.add(Dropout(0.5))

model_1.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    →0, stddev=0.176, seed=None)) )
model_1.add(Dropout(0.5))

model_1.add(Dense(output_dim, activation='softmax'))

model_1.summary()
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 128)	100480

```

-----
dropout_3 (Dropout)          (None, 128)          0
-----
dense_5 (Dense)              (None, 64)           8256
-----
dropout_4 (Dropout)          (None, 64)           0
-----
dense_6 (Dense)              (None, 10)           650
=====
Total params: 109,386
Trainable params: 109,386
Non-trainable params: 0
-----

```

```

[17]: model_1.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 1s 22us/step - loss: 1.2017 -
acc: 0.6102 - val_loss: 0.3453 - val_acc: 0.9066
Epoch 2/20
60000/60000 [=====] - 1s 18us/step - loss: 0.5247 -
acc: 0.8416 - val_loss: 0.2322 - val_acc: 0.9323
Epoch 3/20
60000/60000 [=====] - 1s 17us/step - loss: 0.4000 -
acc: 0.8858 - val_loss: 0.1899 - val_acc: 0.9445
Epoch 4/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3436 -
acc: 0.9016 - val_loss: 0.1686 - val_acc: 0.9523
Epoch 5/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3039 -
acc: 0.9153 - val_loss: 0.1504 - val_acc: 0.9574
Epoch 6/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2777 -
acc: 0.9220 - val_loss: 0.1411 - val_acc: 0.9607
Epoch 7/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2511 -
acc: 0.9294 - val_loss: 0.1360 - val_acc: 0.9617
Epoch 8/20
60000/60000 [=====] - 1s 18us/step - loss: 0.2404 -
acc: 0.9340 - val_loss: 0.1266 - val_acc: 0.9651
Epoch 9/20
60000/60000 [=====] - 1s 18us/step - loss: 0.2240 -

```

```

acc: 0.9375 - val_loss: 0.1191 - val_acc: 0.9672
Epoch 10/20
60000/60000 [=====] - 1s 18us/step - loss: 0.2114 -
acc: 0.9424 - val_loss: 0.1192 - val_acc: 0.9669
Epoch 11/20
60000/60000 [=====] - 1s 18us/step - loss: 0.2034 -
acc: 0.9430 - val_loss: 0.1119 - val_acc: 0.9690
Epoch 12/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1964 -
acc: 0.9446 - val_loss: 0.1155 - val_acc: 0.9685
Epoch 13/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1847 -
acc: 0.9475 - val_loss: 0.1077 - val_acc: 0.9690
Epoch 14/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1807 -
acc: 0.9495 - val_loss: 0.1129 - val_acc: 0.9685
Epoch 15/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1763 -
acc: 0.9503 - val_loss: 0.1066 - val_acc: 0.9713
Epoch 16/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1687 -
acc: 0.9525 - val_loss: 0.1092 - val_acc: 0.9701
Epoch 17/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1618 -
acc: 0.9547 - val_loss: 0.1053 - val_acc: 0.9736
Epoch 18/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1603 -
acc: 0.9538 - val_loss: 0.1090 - val_acc: 0.9725
Epoch 19/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1588 -
acc: 0.9558 - val_loss: 0.1049 - val_acc: 0.9724
Epoch 20/20
60000/60000 [=====] - 1s 18us/step - loss: 0.1526 -
acc: 0.9562 - val_loss: 0.1074 - val_acc: 0.9727

```

```

[18]: score = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↳ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter
    ↳ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10739179229954025

Test accuracy: 0.9727



Model 1. 2 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[19]: # <https://stackoverflow.com/questions/34716454/>
 ↳ where-do-i-call-the-batchnormalization-function-in-keras

```

model_1 = Sequential()

model_1.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None),
    ↳input_shape=(input_dim,)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.2))

model_1.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.2))

model_1.add(Dense(output_dim, activation='softmax'))

model_1.summary()

```

```

-----
Layer (type)                Output Shape              Param #
=====
dense_7 (Dense)              (None, 128)               100480
-----
batch_normalization_3 (Batch Normalization) (None, 128)               512
-----
dropout_5 (Dropout)          (None, 128)               0
-----
dense_8 (Dense)              (None, 64)                8256
-----
batch_normalization_4 (Batch Normalization) (None, 64)               256
-----
dropout_6 (Dropout)          (None, 64)               0
-----
dense_9 (Dense)              (None, 10)                650
=====
Total params: 110,154
Trainable params: 109,770
Non-trainable params: 384
-----

```

```

[20]: model_1.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↳verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 34us/step - loss: 0.5259 -
acc: 0.8396 - val_loss: 0.1926 - val_acc: 0.9411

Epoch 2/20

60000/60000 [=====] - 1s 23us/step - loss: 0.2251 -
acc: 0.9331 - val_loss: 0.1399 - val_acc: 0.9580

Epoch 3/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1740 -
acc: 0.9471 - val_loss: 0.1148 - val_acc: 0.9629

Epoch 4/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1438 -
acc: 0.9562 - val_loss: 0.0969 - val_acc: 0.9704

Epoch 5/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1234 -
acc: 0.9619 - val_loss: 0.0919 - val_acc: 0.9718

Epoch 6/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1093 -
acc: 0.9660 - val_loss: 0.0879 - val_acc: 0.9730

Epoch 7/20

60000/60000 [=====] - 1s 23us/step - loss: 0.1014 -
acc: 0.9682 - val_loss: 0.0784 - val_acc: 0.9754

Epoch 8/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0911 -
acc: 0.9715 - val_loss: 0.0847 - val_acc: 0.9753

Epoch 9/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0820 -
acc: 0.9742 - val_loss: 0.0734 - val_acc: 0.9780

Epoch 10/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0773 -
acc: 0.9757 - val_loss: 0.0716 - val_acc: 0.9774

Epoch 11/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0693 -
acc: 0.9778 - val_loss: 0.0744 - val_acc: 0.9782

Epoch 12/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0663 -
acc: 0.9785 - val_loss: 0.0741 - val_acc: 0.9781

Epoch 13/20

60000/60000 [=====] - 1s 24us/step - loss: 0.0611 -
acc: 0.9802 - val_loss: 0.0675 - val_acc: 0.9800

Epoch 14/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0576 -
acc: 0.9809 - val_loss: 0.0713 - val_acc: 0.9787

Epoch 15/20

60000/60000 [=====] - 1s 24us/step - loss: 0.0547 -
acc: 0.9818 - val_loss: 0.0713 - val_acc: 0.9791

Epoch 16/20

60000/60000 [=====] - 1s 23us/step - loss: 0.0543 -


```

acc: 0.9820 - val_loss: 0.0731 - val_acc: 0.9792
Epoch 17/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0510 -
acc: 0.9832 - val_loss: 0.0665 - val_acc: 0.9804
Epoch 18/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0500 -
acc: 0.9829 - val_loss: 0.0697 - val_acc: 0.9801
Epoch 19/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0473 -
acc: 0.9845 - val_loss: 0.0689 - val_acc: 0.9810
Epoch 20/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0448 -
acc: 0.9852 - val_loss: 0.0674 - val_acc: 0.9807

```

```

[21]: score = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
→number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.06737108307137678
Test accuracy: 0.9807

```



This model is stable after 10 epochs. training further will make the model overfit to train data.

Model 1. 2 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[22]: <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras>

```

model_1 = Sequential()

model_1.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None),
    ↳input_shape=(input_dim,)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.8))

model_1.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_1.add(BatchNormalization())
model_1.add(Dropout(0.8))

model_1.add(Dense(output_dim, activation='softmax'))

model_1.summary()

```

W0715 18:44:06.593636 11376 nn_ops.py:4224] Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

W0715 18:44:06.665421 11376 nn_ops.py:4224] Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 128)	100480
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_7 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 64)	8256
batch_normalization_6 (Batch Normalization)	(None, 64)	256
dropout_8 (Dropout)	(None, 64)	0
dense_12 (Dense)	(None, 10)	650
Total params: 110,154		
Trainable params: 109,770		
Non-trainable params: 384		

```
[23]: model_1.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    ↪ verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 39us/step - loss: 2.1479 -
acc: 0.3531 - val_loss: 0.7555 - val_acc: 0.8517

Epoch 2/20

60000/60000 [=====] - 2s 28us/step - loss: 1.2069 -
acc: 0.5817 - val_loss: 0.5305 - val_acc: 0.8823

Epoch 3/20

60000/60000 [=====] - 2s 28us/step - loss: 0.9829 -
acc: 0.6711 - val_loss: 0.4108 - val_acc: 0.8977

Epoch 4/20

60000/60000 [=====] - 2s 28us/step - loss: 0.8540 -
acc: 0.7230 - val_loss: 0.3469 - val_acc: 0.9060

Epoch 5/20
60000/60000 [=====] - 2s 28us/step - loss: 0.7693 -
acc: 0.7580 - val_loss: 0.3085 - val_acc: 0.9151

Epoch 6/20
60000/60000 [=====] - 2s 28us/step - loss: 0.7085 -
acc: 0.7846 - val_loss: 0.2826 - val_acc: 0.9192

Epoch 7/20
60000/60000 [=====] - 2s 29us/step - loss: 0.6728 -
acc: 0.7994 - val_loss: 0.2663 - val_acc: 0.9245

Epoch 8/20
60000/60000 [=====] - 2s 28us/step - loss: 0.6345 -
acc: 0.8139 - val_loss: 0.2589 - val_acc: 0.9251

Epoch 9/20
60000/60000 [=====] - 2s 28us/step - loss: 0.6115 -
acc: 0.8207 - val_loss: 0.2438 - val_acc: 0.9298

Epoch 10/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5844 -
acc: 0.8289 - val_loss: 0.2397 - val_acc: 0.9298

Epoch 11/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5683 -
acc: 0.8353 - val_loss: 0.2299 - val_acc: 0.9320

Epoch 12/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5589 -
acc: 0.8387 - val_loss: 0.2266 - val_acc: 0.9326

Epoch 13/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5487 -
acc: 0.8437 - val_loss: 0.2190 - val_acc: 0.9345

Epoch 14/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5323 -
acc: 0.8486 - val_loss: 0.2221 - val_acc: 0.9364

Epoch 15/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5266 -
acc: 0.8495 - val_loss: 0.2125 - val_acc: 0.9379

Epoch 16/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5153 -
acc: 0.8545 - val_loss: 0.2090 - val_acc: 0.9400

Epoch 17/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5137 -
acc: 0.8551 - val_loss: 0.2014 - val_acc: 0.9419

Epoch 18/20
60000/60000 [=====] - 2s 28us/step - loss: 0.5073 -
acc: 0.8561 - val_loss: 0.2042 - val_acc: 0.9410

Epoch 19/20
60000/60000 [=====] - 2s 28us/step - loss: 0.4995 -
acc: 0.8579 - val_loss: 0.2001 - val_acc: 0.9438

Epoch 20/20
60000/60000 [=====] - 2s 28us/step - loss: 0.4976 -
acc: 0.8595 - val_loss: 0.1978 - val_acc: 0.9443

```

[24]: score = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→ validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
→ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.19782737234905362

Test accuracy: 0.9443



Due to high dropout rate the model became stable at a high loss than other model and after 10 epoch it is not moving towards the minima

Model 2. 3 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[25]: <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras>

```

model_2 = Sequential()

model_2.add(Dense(512, activation='relu', input_shape=(input_dim,),
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_2.add(BatchNormalization())
model_2.add(Dropout(0.5))

model_2.add(Dense(128, activation='relu',
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.5))

model_2.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    →0, stddev=0.176, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.5))

```

```
model_2.add(Dense(output_dim, activation='softmax'))

model_2.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_13 (Dense)             (None, 512)              401920
-----
batch_normalization_7 (Batch Normalization) (None, 512)              2048
-----
dropout_9 (Dropout)          (None, 512)              0
-----
dense_14 (Dense)             (None, 128)              65664
-----
batch_normalization_8 (Batch Normalization) (None, 128)              512
-----
dropout_10 (Dropout)         (None, 128)              0
-----
dense_15 (Dense)             (None, 64)               8256
-----
batch_normalization_9 (Batch Normalization) (None, 64)              256
-----
dropout_11 (Dropout)         (None, 64)              0
-----
dense_16 (Dense)             (None, 10)               650
=====
Total params: 479,306
Trainable params: 477,898
Non-trainable params: 1,408
-----
```

```
[26]: model_2.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 98us/step - loss: 0.8009 -
acc: 0.7527 - val_loss: 0.2047 - val_acc: 0.9364

Epoch 2/20

60000/60000 [=====] - 5s 81us/step - loss: 0.3331 -
acc: 0.9038 - val_loss: 0.1434 - val_acc: 0.9528

Epoch 3/20

60000/60000 [=====] - 5s 82us/step - loss: 0.2573 -
acc: 0.9271 - val_loss: 0.1162 - val_acc: 0.9629
Epoch 4/20
60000/60000 [=====] - 5s 82us/step - loss: 0.2104 -
acc: 0.9410 - val_loss: 0.0986 - val_acc: 0.9690
Epoch 5/20
60000/60000 [=====] - 5s 81us/step - loss: 0.1828 -
acc: 0.9490 - val_loss: 0.0913 - val_acc: 0.9751
Epoch 6/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1611 -
acc: 0.9559 - val_loss: 0.0885 - val_acc: 0.9738
Epoch 7/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1474 -
acc: 0.9576 - val_loss: 0.0812 - val_acc: 0.9766
Epoch 8/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1373 -
acc: 0.9615 - val_loss: 0.0796 - val_acc: 0.9756
Epoch 9/20
60000/60000 [=====] - 5s 81us/step - loss: 0.1270 -
acc: 0.9643 - val_loss: 0.0718 - val_acc: 0.9788
Epoch 10/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1222 -
acc: 0.9661 - val_loss: 0.0678 - val_acc: 0.9805
Epoch 11/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1143 -
acc: 0.9681 - val_loss: 0.0701 - val_acc: 0.9795
Epoch 12/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1080 -
acc: 0.9700 - val_loss: 0.0665 - val_acc: 0.9816
Epoch 13/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1032 -
acc: 0.9711 - val_loss: 0.0684 - val_acc: 0.9804
Epoch 14/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0951 -
acc: 0.9730 - val_loss: 0.0657 - val_acc: 0.9819
Epoch 15/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0929 -
acc: 0.9742 - val_loss: 0.0633 - val_acc: 0.9828
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0878 -
acc: 0.9751 - val_loss: 0.0635 - val_acc: 0.9835
Epoch 17/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0833 -
acc: 0.9767 - val_loss: 0.0616 - val_acc: 0.9836
Epoch 18/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0812 -
acc: 0.9769 - val_loss: 0.0640 - val_acc: 0.9819
Epoch 19/20


```

60000/60000 [=====] - 5s 82us/step - loss: 0.0797 -
acc: 0.9778 - val_loss: 0.0614 - val_acc: 0.9819
Epoch 20/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0755 -
acc: 0.9777 - val_loss: 0.0632 - val_acc: 0.9823

```

```

[27]: score = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
→ number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.06315016652900377
Test accuracy: 0.9823

```



Model 2. 3 hidden layer, ReLU, Dropout and AdamOptimizer

[28]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras>

```

model_2 = Sequential()

model_2.add(Dense(512, activation='relu', input_shape=(input_dim,),
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_2.add(Dropout(0.5))

model_2.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_2.add(Dropout(0.5))

model_2.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_2.add(Dropout(0.5))

model_2.add(Dense(output_dim, activation='softmax'))

model_2.summary()

```

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
dropout_12 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
dropout_13 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 64)	8256
dropout_14 (Dropout)	(None, 64)	0
dense_20 (Dense)	(None, 10)	650
Total params: 476,490		
Trainable params: 476,490		
Non-trainable params: 0		

```
[29]: model_2.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 84us/step - loss: 1.1420 -
acc: 0.6327 - val_loss: 0.2604 - val_acc: 0.9287
Epoch 2/20
60000/60000 [=====] - 4s 72us/step - loss: 0.4306 -
acc: 0.8802 - val_loss: 0.1797 - val_acc: 0.9508
Epoch 3/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3154 -
acc: 0.9168 - val_loss: 0.1521 - val_acc: 0.9590
Epoch 4/20
60000/60000 [=====] - 4s 71us/step - loss: 0.2582 -
acc: 0.9323 - val_loss: 0.1386 - val_acc: 0.9640
Epoch 5/20
60000/60000 [=====] - 4s 72us/step - loss: 0.2209 -
acc: 0.9430 - val_loss: 0.1261 - val_acc: 0.9664
Epoch 6/20
60000/60000 [=====] - 4s 72us/step - loss: 0.2000 -
acc: 0.9481 - val_loss: 0.1072 - val_acc: 0.9723
```

```

Epoch 7/20
60000/60000 [=====] - ETA: 0s - loss: 0.1823 - acc:
0.952 - 4s 72us/step - loss: 0.1826 - acc: 0.9527 - val_loss: 0.0997 - val_acc:
0.9737
Epoch 8/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1682 -
acc: 0.9566 - val_loss: 0.1014 - val_acc: 0.9734
Epoch 9/20
60000/60000 [=====] - 4s 73us/step - loss: 0.1548 -
acc: 0.9606 - val_loss: 0.0945 - val_acc: 0.9763
Epoch 10/20
60000/60000 [=====] - 4s 73us/step - loss: 0.1437 -
acc: 0.9630 - val_loss: 0.0952 - val_acc: 0.9754
Epoch 11/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1382 -
acc: 0.9647 - val_loss: 0.0938 - val_acc: 0.9760
Epoch 12/20
60000/60000 [=====] - 5s 75us/step - loss: 0.1300 -
acc: 0.9667 - val_loss: 0.0891 - val_acc: 0.9778
Epoch 13/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1217 -
acc: 0.9685 - val_loss: 0.0863 - val_acc: 0.9786
Epoch 14/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1162 -
acc: 0.9703 - val_loss: 0.0892 - val_acc: 0.9780
Epoch 15/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1105 -
acc: 0.9711 - val_loss: 0.0904 - val_acc: 0.9790
Epoch 16/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1038 -
acc: 0.9720 - val_loss: 0.0845 - val_acc: 0.9794
Epoch 17/20
60000/60000 [=====] - 4s 73us/step - loss: 0.1040 -
acc: 0.9728 - val_loss: 0.0818 - val_acc: 0.9784
Epoch 18/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1001 -
acc: 0.9741 - val_loss: 0.0846 - val_acc: 0.9800
Epoch 19/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0953 -
acc: 0.9750 - val_loss: 0.0802 - val_acc: 0.9796
Epoch 20/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0877 -
acc: 0.9768 - val_loss: 0.0831 - val_acc: 0.9807

```

```

[30]: score = model_2.evaluate(X_test, Y_test, verbose=0)
      print('Test score:', score[0])
      print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    →epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
    →validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    →number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08312617581425139

Test accuracy: 0.9807



Model 3. 3 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[31]: # <https://stackoverflow.com/questions/34716454/>

→where-do-i-call-the-batchnormalization-function-in-keras

```
model_2 = Sequential()

model_2.add(Dense(512, activation='relu', input_shape=(input_dim,),
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_2.add(BatchNormalization())
model_2.add(Dropout(0.2))

model_2.add(Dense(128, activation='relu',
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.2))

model_2.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    →0, stddev=0.176, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.2))

model_2.add(Dense(output_dim, activation='softmax'))
```

```
model_2.summary()
```

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 512)	401920
batch_normalization_10 (Batch Normalization)	(None, 512)	2048
dropout_15 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 128)	65664
batch_normalization_11 (Batch Normalization)	(None, 128)	512
dropout_16 (Dropout)	(None, 128)	0
dense_23 (Dense)	(None, 64)	8256
batch_normalization_12 (Batch Normalization)	(None, 64)	256
dropout_17 (Dropout)	(None, 64)	0
dense_24 (Dense)	(None, 10)	650

Total params: 479,306
Trainable params: 477,898
Non-trainable params: 1,408

```
[32]: model_2.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    ↪ verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 104us/step - loss: 0.3539 -
acc: 0.8956 - val_loss: 0.1141 - val_acc: 0.9654

Epoch 2/20

60000/60000 [=====] - 5s 84us/step - loss: 0.1448 -
acc: 0.9573 - val_loss: 0.0891 - val_acc: 0.9732

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1093 -
acc: 0.9661 - val_loss: 0.0794 - val_acc: 0.9747

Epoch 4/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0881 -
acc: 0.9730 - val_loss: 0.0681 - val_acc: 0.9787

Epoch 5/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0756 -
acc: 0.9760 - val_loss: 0.0728 - val_acc: 0.9780

Epoch 6/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0660 -
acc: 0.9797 - val_loss: 0.0626 - val_acc: 0.9802

Epoch 7/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0566 -
acc: 0.9823 - val_loss: 0.0657 - val_acc: 0.9810

Epoch 8/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0501 -
acc: 0.9838 - val_loss: 0.0640 - val_acc: 0.9812

Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0472 -
acc: 0.9847 - val_loss: 0.0719 - val_acc: 0.9793

Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0461 -
acc: 0.9853 - val_loss: 0.0599 - val_acc: 0.9829

Epoch 11/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0397 -
acc: 0.9870 - val_loss: 0.0600 - val_acc: 0.9814

Epoch 12/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0387 -
acc: 0.9870 - val_loss: 0.0537 - val_acc: 0.9841

Epoch 13/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0356 -
acc: 0.9885 - val_loss: 0.0624 - val_acc: 0.9825

Epoch 14/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0347 -
acc: 0.9884 - val_loss: 0.0557 - val_acc: 0.9847

Epoch 15/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0316 -
acc: 0.9898 - val_loss: 0.0615 - val_acc: 0.9836

Epoch 16/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0274 -
acc: 0.9906 - val_loss: 0.0664 - val_acc: 0.9829

Epoch 17/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0282 -
acc: 0.9906 - val_loss: 0.0577 - val_acc: 0.9841

Epoch 18/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0262 -
acc: 0.9908 - val_loss: 0.0694 - val_acc: 0.9828

Epoch 19/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0258 -
acc: 0.9914 - val_loss: 0.0618 - val_acc: 0.9830

Epoch 20/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0243 -
acc: 0.9918 - val_loss: 0.0581 - val_acc: 0.9854

```
[33]: score = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#     → epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter
#     → validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
#     → number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.058136306594208875

Test accuracy: 0.9854



Model 3. 3 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[34]: # <https://stackoverflow.com/questions/34716454/>

→where-do-i-call-the-batchnormalization-function-in-keras

```
model_2 = Sequential()

model_2.add(Dense(512, activation='relu', input_shape=(input_dim,),
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_2.add(BatchNormalization())
model_2.add(Dropout(0.8))

model_2.add(Dense(128, activation='relu',
    →kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.8))

model_2.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    →0, stddev=0.176, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.8))

model_2.add(Dense(output_dim, activation='softmax'))
```

```
model_2.summary()
```

W0715 18:49:34.213879 11376 nn_ops.py:4224] Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

W0715 18:49:34.288651 11376 nn_ops.py:4224] Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

W0715 18:49:34.363479 11376 nn_ops.py:4224] Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 512)	401920
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dropout_18 (Dropout)	(None, 512)	0
dense_26 (Dense)	(None, 128)	65664
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_27 (Dense)	(None, 64)	8256
batch_normalization_15 (Batch Normalization)	(None, 64)	256
dropout_20 (Dropout)	(None, 64)	0
dense_28 (Dense)	(None, 10)	650

Total params: 479,306
Trainable params: 477,898
Non-trainable params: 1,408

```
[35]: model_2.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    ↪ verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 6s 105us/step - loss: 2.4937 -
acc: 0.2566 - val_loss: 1.1451 - val_acc: 0.7325

Epoch 2/20
60000/60000 [=====] - 5s 85us/step - loss: 1.4336 -
acc: 0.4918 - val_loss: 0.7489 - val_acc: 0.8238

Epoch 3/20
60000/60000 [=====] - 5s 85us/step - loss: 1.1254 -
acc: 0.6133 - val_loss: 0.5330 - val_acc: 0.8762

Epoch 4/20
60000/60000 [=====] - 5s 86us/step - loss: 0.9454 -
acc: 0.6819 - val_loss: 0.4127 - val_acc: 0.9009

Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.8193 -
acc: 0.7359 - val_loss: 0.3276 - val_acc: 0.9199

Epoch 6/20
60000/60000 [=====] - 5s 85us/step - loss: 0.7250 -
acc: 0.7723 - val_loss: 0.2828 - val_acc: 0.9273

Epoch 7/20
60000/60000 [=====] - 5s 85us/step - loss: 0.6651 -
acc: 0.7997 - val_loss: 0.2425 - val_acc: 0.9375

Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.6095 -
acc: 0.8204 - val_loss: 0.2189 - val_acc: 0.9425

Epoch 9/20
60000/60000 [=====] - 5s 86us/step - loss: 0.5651 -
acc: 0.8388 - val_loss: 0.1981 - val_acc: 0.9469

Epoch 10/20
60000/60000 [=====] - 5s 86us/step - loss: 0.5232 -
acc: 0.8530 - val_loss: 0.1911 - val_acc: 0.9496

Epoch 11/20
60000/60000 [=====] - 5s 86us/step - loss: 0.5006 -
acc: 0.8616 - val_loss: 0.1793 - val_acc: 0.9520

Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4841 -
acc: 0.8678 - val_loss: 0.1732 - val_acc: 0.9540

Epoch 13/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4584 -
acc: 0.8784 - val_loss: 0.1608 - val_acc: 0.9574

Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.4446 -
acc: 0.8804 - val_loss: 0.1570 - val_acc: 0.9583

Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4335 -
acc: 0.8841 - val_loss: 0.1536 - val_acc: 0.9584

Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4205 -

```

acc: 0.8886 - val_loss: 0.1567 - val_acc: 0.9588
Epoch 17/20
60000/60000 [=====] - 5s 87us/step - loss: 0.4047 -
acc: 0.8940 - val_loss: 0.1533 - val_acc: 0.9594
Epoch 18/20
60000/60000 [=====] - 5s 86us/step - loss: 0.3985 -
acc: 0.8953 - val_loss: 0.1441 - val_acc: 0.9626
Epoch 19/20
60000/60000 [=====] - 5s 86us/step - loss: 0.3938 -
acc: 0.8989 - val_loss: 0.1454 - val_acc: 0.9637
Epoch 20/20
60000/60000 [=====] - 5s 85us/step - loss: 0.3815 -
acc: 0.8997 - val_loss: 0.1369 - val_acc: 0.9648

```

```

[36]: score = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
→ number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.1369463347967714
Test accuracy: 0.9648

```



Model 3. 5 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[37]: <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras>

```

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model_3 = Sequential()

model_3.add(Dense(512, activation='relu', input_shape=(input_dim,),
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))

model_3.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))

```

```

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
→0, stddev=0.176, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))

model_3.add(Dense(128, activation='relu',
→kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()

```

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 512)	401920
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dropout_21 (Dropout)	(None, 512)	0
dense_30 (Dense)	(None, 128)	65664
batch_normalization_17 (Batch Normalization)	(None, 128)	512
dropout_22 (Dropout)	(None, 128)	0
dense_31 (Dense)	(None, 64)	8256
batch_normalization_18 (Batch Normalization)	(None, 64)	256
dropout_23 (Dropout)	(None, 64)	0
dense_32 (Dense)	(None, 64)	4160
batch_normalization_19 (Batch Normalization)	(None, 64)	256
dropout_24 (Dropout)	(None, 64)	0
dense_33 (Dense)	(None, 128)	8320
batch_normalization_20 (Batch Normalization)	(None, 128)	512
dropout_25 (Dropout)	(None, 128)	0

```

-----
dense_34 (Dense)                (None, 10)                1290
=====
Total params: 493,194
Trainable params: 491,402
Non-trainable params: 1,792
-----

```

```

[38]: model_3.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 128us/step - loss: 1.6475 -
acc: 0.4693 - val_loss: 0.4033 - val_acc: 0.8867
Epoch 2/20
60000/60000 [=====] - 6s 96us/step - loss: 0.6705 -
acc: 0.7905 - val_loss: 0.2357 - val_acc: 0.9349
Epoch 3/20
60000/60000 [=====] - 6s 97us/step - loss: 0.4587 -
acc: 0.8699 - val_loss: 0.1786 - val_acc: 0.9511
Epoch 4/20
60000/60000 [=====] - 6s 96us/step - loss: 0.3607 -
acc: 0.9007 - val_loss: 0.1519 - val_acc: 0.9588
Epoch 5/20
60000/60000 [=====] - 6s 97us/step - loss: 0.3025 -
acc: 0.9191 - val_loss: 0.1399 - val_acc: 0.9642
Epoch 6/20
60000/60000 [=====] - 6s 98us/step - loss: 0.2703 -
acc: 0.9301 - val_loss: 0.1225 - val_acc: 0.9695
Epoch 7/20
60000/60000 [=====] - 6s 99us/step - loss: 0.2407 -
acc: 0.9386 - val_loss: 0.1176 - val_acc: 0.9689
Epoch 8/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2219 -
acc: 0.9451 - val_loss: 0.1185 - val_acc: 0.9700
Epoch 9/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2048 -
acc: 0.9494 - val_loss: 0.1026 - val_acc: 0.9740
Epoch 10/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1918 -
acc: 0.9522 - val_loss: 0.1019 - val_acc: 0.9751
Epoch 11/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1762 -

```



```

acc: 0.9569 - val_loss: 0.0951 - val_acc: 0.9767
Epoch 12/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1703 -
acc: 0.9581 - val_loss: 0.0930 - val_acc: 0.9769
Epoch 13/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1579 -
acc: 0.9616 - val_loss: 0.1007 - val_acc: 0.9762
Epoch 14/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1539 -
acc: 0.9613 - val_loss: 0.0892 - val_acc: 0.9781
Epoch 15/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1483 -
acc: 0.9635 - val_loss: 0.0816 - val_acc: 0.9803
Epoch 16/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1425 -
acc: 0.9636 - val_loss: 0.0866 - val_acc: 0.9797
Epoch 17/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1366 -
acc: 0.9661 - val_loss: 0.0857 - val_acc: 0.9795
Epoch 18/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1309 -
acc: 0.9678 - val_loss: 0.0832 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1290 -
acc: 0.9681 - val_loss: 0.0840 - val_acc: 0.9812
Epoch 20/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1218 -
acc: 0.9696 - val_loss: 0.0797 - val_acc: 0.9823

```

```

[39]: score = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter
→validation_data
# val_loss : validation loss

```

```

# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07966478181201965

Test accuracy: 0.9823



Model 3. 5 hidden layer, ReLU, Dropout and AdamOptimizer

```

[40]: # https://stackoverflow.com/questions/34716454/
    ↳ where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model_3 = Sequential()

model_3.add(Dense(512, activation='relu', input_shape=(input_dim,),
    ↳ kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))

```

```

model_3.add(Dropout(0.5))

model_3.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(Dropout(0.5))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_3.add(Dropout(0.5))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_3.add(Dropout(0.5))

model_3.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(Dropout(0.5))

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()

```

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 128)	65664
dropout_27 (Dropout)	(None, 128)	0
dense_37 (Dense)	(None, 64)	8256
dropout_28 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 64)	4160
dropout_29 (Dropout)	(None, 64)	0
dense_39 (Dense)	(None, 128)	8320
dropout_30 (Dropout)	(None, 128)	0
dense_40 (Dense)	(None, 10)	1290

```
=====
Total params: 489,610
Trainable params: 489,610
Non-trainable params: 0
-----
```

```
[41]: model_3.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪metrics=['accuracy'])  
  
history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    ↪verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 103us/step - loss: 2.1719 -  
acc: 0.2328 - val_loss: 1.4112 - val_acc: 0.4797
Epoch 2/20
60000/60000 [=====] - 5s 82us/step - loss: 1.2927 -  
acc: 0.5287 - val_loss: 0.8277 - val_acc: 0.6854
Epoch 3/20
60000/60000 [=====] - 5s 82us/step - loss: 0.9126 -  
acc: 0.6637 - val_loss: 0.5805 - val_acc: 0.8139
Epoch 4/20
60000/60000 [=====] - 5s 82us/step - loss: 0.7038 -  
acc: 0.7515 - val_loss: 0.4486 - val_acc: 0.8323
Epoch 5/20
60000/60000 [=====] - 5s 82us/step - loss: 0.5840 -  
acc: 0.7958 - val_loss: 0.3832 - val_acc: 0.8513
Epoch 6/20
60000/60000 [=====] - 5s 82us/step - loss: 0.5204 -  
acc: 0.8157 - val_loss: 0.3473 - val_acc: 0.8589
Epoch 7/20
60000/60000 [=====] - 5s 83us/step - loss: 0.4657 -  
acc: 0.8338 - val_loss: 0.3145 - val_acc: 0.8690
Epoch 8/20
60000/60000 [=====] - 5s 86us/step - loss: 0.4365 -  
acc: 0.8521 - val_loss: 0.2935 - val_acc: 0.9065
Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.3839 -  
acc: 0.8963 - val_loss: 0.2154 - val_acc: 0.9530
Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.3319 -  
acc: 0.9199 - val_loss: 0.1913 - val_acc: 0.9574
Epoch 11/20
60000/60000 [=====] - 5s 82us/step - loss: 0.3073 -  
acc: 0.9251 - val_loss: 0.1842 - val_acc: 0.9573
Epoch 12/20
```

```

60000/60000 [=====] - 5s 83us/step - loss: 0.2882 -
acc: 0.9305 - val_loss: 0.1751 - val_acc: 0.9607
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.2737 -
acc: 0.9349 - val_loss: 0.1632 - val_acc: 0.9627
Epoch 14/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2538 -
acc: 0.9404 - val_loss: 0.1705 - val_acc: 0.9638
Epoch 15/20
60000/60000 [=====] - 5s 83us/step - loss: 0.2426 -
acc: 0.9431 - val_loss: 0.1562 - val_acc: 0.9662
Epoch 16/20
60000/60000 [=====] - 5s 83us/step - loss: 0.2400 -
acc: 0.9442 - val_loss: 0.1511 - val_acc: 0.9655
Epoch 17/20
60000/60000 [=====] - 5s 82us/step - loss: 0.2259 -
acc: 0.9467 - val_loss: 0.1447 - val_acc: 0.9681
Epoch 18/20
60000/60000 [=====] - 5s 82us/step - loss: 0.2184 -
acc: 0.9498 - val_loss: 0.1446 - val_acc: 0.9691
Epoch 19/20
60000/60000 [=====] - 5s 83us/step - loss: 0.2153 -
acc: 0.9485 - val_loss: 0.1456 - val_acc: 0.9665
Epoch 20/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2050 -
acc: 0.9515 - val_loss: 0.1340 - val_acc: 0.9692

```

```

[42]: score = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→validation_data
# val_loss : validation loss
# val_acc : validation accuracy


```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.13398695126250387

Test accuracy: 0.9692



Model 3. 5 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

[43]: # <https://stackoverflow.com/questions/34716454/>
 ↳where-do-i-call-the-batchnormalization-function-in-keras

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model_3 = Sequential()

model_3.add(Dense(512, activation='relu', input_shape=(input_dim,),
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_3.add(BatchNormalization())
```

```

model_3.add(Dropout(0.2))

model_3.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.2))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.2))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
    ↳0, stddev=0.176, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.2))

model_3.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.2))

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()

```

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 512)	401920
batch_normalization_21 (Batch Normalization)	(None, 512)	2048
dropout_31 (Dropout)	(None, 512)	0
dense_42 (Dense)	(None, 128)	65664
batch_normalization_22 (Batch Normalization)	(None, 128)	512
dropout_32 (Dropout)	(None, 128)	0
dense_43 (Dense)	(None, 64)	8256
batch_normalization_23 (Batch Normalization)	(None, 64)	256
dropout_33 (Dropout)	(None, 64)	0

```

-----
dense_44 (Dense)                (None, 64)                4160
-----
batch_normalization_24 (Batch Normalization) (None, 64)                256
-----
dropout_34 (Dropout)            (None, 64)                0
-----
dense_45 (Dense)                (None, 128)               8320
-----
batch_normalization_25 (Batch Normalization) (None, 128)              512
-----
dropout_35 (Dropout)            (None, 128)               0
-----
dense_46 (Dense)                (None, 10)               1290
=====
Total params: 493,194
Trainable params: 491,402
Non-trainable params: 1,792
-----

```

```

[44]: model_3.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 139us/step - loss: 0.5371 -
acc: 0.8352 - val_loss: 0.1464 - val_acc: 0.9552

Epoch 2/20

60000/60000 [=====] - 6s 104us/step - loss: 0.2128 -
acc: 0.9385 - val_loss: 0.1091 - val_acc: 0.9672

Epoch 3/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1594 -
acc: 0.9534 - val_loss: 0.0922 - val_acc: 0.9725

Epoch 4/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1300 -
acc: 0.9622 - val_loss: 0.0867 - val_acc: 0.9734

Epoch 5/20

60000/60000 [=====] - 6s 106us/step - loss: 0.1093 -
acc: 0.9679 - val_loss: 0.0753 - val_acc: 0.9784

Epoch 6/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0980 -
acc: 0.9712 - val_loss: 0.0803 - val_acc: 0.9770

Epoch 7/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0893 -


```

acc: 0.9731 - val_loss: 0.0781 - val_acc: 0.9778
Epoch 8/20
60000/60000 [=====] - 7s 111us/step - loss: 0.0797 -
acc: 0.9765 - val_loss: 0.0743 - val_acc: 0.9797
Epoch 9/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0732 -
acc: 0.9779 - val_loss: 0.0685 - val_acc: 0.9815
Epoch 10/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0677 -
acc: 0.9803 - val_loss: 0.0658 - val_acc: 0.9817
Epoch 11/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0612 -
acc: 0.9820 - val_loss: 0.0652 - val_acc: 0.9815
Epoch 12/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0586 -
acc: 0.9824 - val_loss: 0.0692 - val_acc: 0.9810
Epoch 13/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0542 -
acc: 0.9837 - val_loss: 0.0694 - val_acc: 0.9811
Epoch 14/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0477 -
acc: 0.9855 - val_loss: 0.0691 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0503 -
acc: 0.9850 - val_loss: 0.0828 - val_acc: 0.9785
Epoch 16/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0454 -
acc: 0.9861 - val_loss: 0.0699 - val_acc: 0.9813
Epoch 17/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0429 -
acc: 0.9868 - val_loss: 0.0683 - val_acc: 0.9830
Epoch 18/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0417 -
acc: 0.9870 - val_loss: 0.0746 - val_acc: 0.9804
Epoch 19/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0410 -
acc: 0.9874 - val_loss: 0.0616 - val_acc: 0.9832
Epoch 20/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0393 -
acc: 0.9880 - val_loss: 0.0631 - val_acc: 0.9829

```

```

[45]: score = model_3.evaluate(X_test, Y_test, verbose=0)
      print('Test score:', score[0])
      print('Test accuracy:', score[1])

      fig,ax = plt.subplots(1,1)
      ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    →epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
    →validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    →number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06305098229350987

Test accuracy: 0.9829



Model 3. 5 hidden layer, ReLU, Dropout, BatchNormalization and AdamOptimizer

```
[46]: # https://stackoverflow.com/questions/34716454/
      ↪ where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model_3 = Sequential()

model_3.add(Dense(512, activation='relu', input_shape=(input_dim, ),
      ↪ kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.8))

model_3.add(Dense(128, activation='relu',
      ↪ kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.8))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
      ↪ 0, stddev=0.176, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.8))

model_3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.
      ↪ 0, stddev=0.176, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.8))

model_3.add(Dense(128, activation='relu',
      ↪ kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_3.add(BatchNormalization())
model_3.add(Dropout(0.8))

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()
```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 512)	401920
batch_normalization_26 (Batc	(None, 512)	2048
dropout_36 (Dropout)	(None, 512)	0

```

-----
dense_48 (Dense)                (None, 128)                65664
-----
batch_normalization_27 (Batch Normalization) (None, 128)                512
-----
dropout_37 (Dropout)            (None, 128)                0
-----
dense_49 (Dense)                (None, 64)                 8256
-----
batch_normalization_28 (Batch Normalization) (None, 64)                256
-----
dropout_38 (Dropout)            (None, 64)                0
-----
dense_50 (Dense)                (None, 64)                4160
-----
batch_normalization_29 (Batch Normalization) (None, 64)                256
-----
dropout_39 (Dropout)            (None, 64)                0
-----
dense_51 (Dense)                (None, 128)               8320
-----
batch_normalization_30 (Batch Normalization) (None, 128)               512
-----
dropout_40 (Dropout)            (None, 128)                0
-----
dense_52 (Dense)                (None, 10)                1290
=====
Total params: 493,194
Trainable params: 491,402
Non-trainable params: 1,792
-----

```

```

[47]: model_3.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 144us/step - loss: 3.0380 -
acc: 0.1137 - val_loss: 2.3314 - val_acc: 0.1145

Epoch 2/20

60000/60000 [=====] - 6s 106us/step - loss: 2.2017 -
acc: 0.1747 - val_loss: 2.2887 - val_acc: 0.1366

Epoch 3/20

60000/60000 [=====] - 6s 103us/step - loss: 2.0982 -

acc: 0.2019 - val_loss: 2.2321 - val_acc: 0.1283
Epoch 4/20
60000/60000 [=====] - 6s 103us/step - loss: 2.0141 -
acc: 0.2138 - val_loss: 2.1096 - val_acc: 0.1461
Epoch 5/20
60000/60000 [=====] - 7s 112us/step - loss: 1.9522 -
acc: 0.2232 - val_loss: 1.9815 - val_acc: 0.1942
Epoch 6/20
60000/60000 [=====] - 6s 104us/step - loss: 1.8904 -
acc: 0.2424 - val_loss: 1.8959 - val_acc: 0.2227
Epoch 7/20
60000/60000 [=====] - 6s 106us/step - loss: 1.8227 -
acc: 0.2633 - val_loss: 1.8004 - val_acc: 0.2619
Epoch 8/20
60000/60000 [=====] - 6s 105us/step - loss: 1.7593 -
acc: 0.2829 - val_loss: 1.6761 - val_acc: 0.3008
Epoch 9/20
60000/60000 [=====] - 6s 105us/step - loss: 1.6793 -
acc: 0.3099 - val_loss: 1.4377 - val_acc: 0.4360
Epoch 10/20
60000/60000 [=====] - 6s 104us/step - loss: 1.6044 -
acc: 0.3410 - val_loss: 1.2655 - val_acc: 0.4952
Epoch 11/20
60000/60000 [=====] - 6s 103us/step - loss: 1.5093 -
acc: 0.3793 - val_loss: 1.1853 - val_acc: 0.4824
Epoch 12/20
60000/60000 [=====] - 6s 104us/step - loss: 1.4331 -
acc: 0.4036 - val_loss: 1.1400 - val_acc: 0.5061
Epoch 13/20
60000/60000 [=====] - 7s 111us/step - loss: 1.3994 -
acc: 0.4158 - val_loss: 1.0967 - val_acc: 0.5307
Epoch 14/20
60000/60000 [=====] - 7s 109us/step - loss: 1.3628 -
acc: 0.4262 - val_loss: 1.0733 - val_acc: 0.5401
Epoch 15/20
60000/60000 [=====] - 7s 110us/step - loss: 1.3399 -
acc: 0.4352 - val_loss: 1.0520 - val_acc: 0.5453
Epoch 16/20
60000/60000 [=====] - 6s 106us/step - loss: 1.3246 -
acc: 0.4398 - val_loss: 1.0232 - val_acc: 0.6262
Epoch 17/20
60000/60000 [=====] - 7s 108us/step - loss: 1.3028 -
acc: 0.4462 - val_loss: 1.0212 - val_acc: 0.5690
Epoch 18/20
60000/60000 [=====] - 6s 108us/step - loss: 1.2857 -
acc: 0.4552 - val_loss: 1.0059 - val_acc: 0.5607
Epoch 19/20
60000/60000 [=====] - 7s 110us/step - loss: 1.2704 -

```
acc: 0.4587 - val_loss: 0.9887 - val_acc: 0.5915
Epoch 20/20
60000/60000 [=====] - 7s 110us/step - loss: 1.2616 -
acc: 0.4644 - val_loss: 0.9821 - val_acc: 0.5892
```

```
[48]: score = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    →epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
    →validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    →number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.9821138044357299
Test accuracy: 0.5892
```



```
[49]: from prettytable import PrettyTable
y = PrettyTable()
y.field_names = ["Model 1", "Model 2", "Model 3"]
y.add_row(["128 - 64", "512 - 128 - 64", "512 - 128 - 64 - 64 - 128"])
y.add_row(["2 layer NN", "3 layer NN", "5 layer NN"])
print(y)

x = PrettyTable()
x.field_names = ["Model", "Dropout", "Batch Normalization", "optimizer", "Train_
→Accuracy", "Test Accuracy"]
x.add_row(["Model 1.1", "0.5", "Yes", "adam", 0.959, 0.976])
x.add_row(["Model 1.2", "0.5", "No", "adam", 0.956, 0.979])
x.add_row(["Model 1.3", "0.2", "Yes", "adam", 0.985, 0.979])
x.add_row(["Model 1.4", "0.8", "Yes", "adam", 0.859, 0.945])

x.add_row(["-----", "----", "----", "----", "----", "----"])

x.add_row(["Model 2.1", "0.5", "Yes", "adam", 0.978, 0.983])
x.add_row(["Model 2.2", "0.5", "No", "adam", 0.975, 0.978])
x.add_row(["Model 2.3", "0.2", "Yes", "adam", 0.992, 0.981])
x.add_row(["Model 2.4", "0.8", "Yes", "adam", 0.902, 0.965])

x.add_row(["-----", "----", "----", "----", "----", "----"])
```

```

x.add_row(["Model 3.1", "0.5", "Yes", "adam",0.971,0.979])
x.add_row(["Model 3.2", "0.5" , "No", "adam",0.951,0.967])
x.add_row(["Model 3.3", "0.2", "Yes", "adam",0.989,0.984])
x.add_row(["Model 3.4", "0.8", "Yes", "adam",0.406,0.502])

x.border=True
print(x)

```

Model 1		Model 2		Model 3	
128 - 64		512 - 128 - 64		512 - 128 - 64 - 64 - 128	
2 layer NN		3 layer NN		5 layer NN	
Model	Dropout	Batch Normalization	optimizer	Train Accuracy	Test Accuracy
Model 1.1	0.5	Yes	adam	0.959	0.976
Model 1.2	0.5	No	adam	0.956	0.979
Model 1.3	0.2	Yes	adam	0.985	0.979
Model 1.4	0.8	Yes	adam	0.859	0.945
Model 2.1	0.5	Yes	adam	0.978	0.983
Model 2.2	0.5	No	adam	0.975	0.978
Model 2.3	0.2	Yes	adam	0.992	0.981
Model 2.4	0.8	Yes	adam	0.902	0.965
Model 3.1	0.5	Yes	adam	0.971	0.979
Model 3.2	0.5	No	adam	0.951	0.967
Model 3.3	0.2	Yes	adam	0.989	0.984
Model 3.4	0.8	Yes	adam	0.406	0.502

0.502 |
+-----+-----+-----+-----+-----+-----+
-----+

Best model is Model 3.3 as it has both train and test score as 0.989 and 0.984 respectively

Drop out of value 0.8 hurts the model. 0.5 dropout is reasonable.

From result 1.2, 2.2, 3.2 its pretty clear that batch Normalization is needed in later stages of deep neural network, especially having more than 3 hidden layers. It is always good to do batch Normalization.

[]: