

prabhudayala@gmail.com_25

October 30, 2019

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve

```

C:\Users\user\Anaconda3\lib\site-packages\sklearn\externals\six.py:31:
DeprecationWarning: The module is deprecated in version 0.21 and will be removed
in version 0.23 since we've dropped support for Python 2.7. Please rely on the
official version of six (<https://pypi.org/project/six/>).
"(https://pypi.org/project/six/).", DeprecationWarning)

4. Machine Learning Models

4.1 Reading data from file

```

[2]: X_train = pd.read_csv('final_features_tfidf_train.csv')
      X_test = pd.read_csv('final_features_tfidf_test.csv')

[3]: y_train = pd.read_csv('final_features_tfidf_train_label.csv',header=None)
      y_test = pd.read_csv('final_features_tfidf_test_label.csv',header=None)
      y_train = y_train[1]
      y_train = y_train.values.reshape((y_train.shape[0],1))
      print(y_train.shape)
      y_test = y_test[1]
      y_test = y_test.values.reshape((y_test.shape[0],1))
      print(y_test.shape)

```

```

(140000, 1)
(60000, 1)

```

```

[4]: X_train.drop(['Unnamed: 0', 'id','is_duplicate'], axis=1, inplace=True)
      X_test.drop(['Unnamed: 0', 'id','is_duplicate'], axis=1, inplace=True)

```

```

[5]: X_train.head()

```

```

[5]:   0_x  1_x  2_x  3_x  4_x  5_x  6_x  7_x  8_x  9_x  ...  freq_qid2  q1len  \
0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...         2      40
1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...         1      50
2  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...        12      37
3  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...         2      60
4  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...         1      59

      q2len  q1_n_words  q2_n_words  word_Common  word_Total  word_share  \
0         33          9           7           6.0          16.0       0.375000

```

| | | | | | | |
|---|-----|----|----|-----|------|----------|
| 1 | 35 | 9 | 8 | 6.0 | 16.0 | 0.375000 |
| 2 | 48 | 7 | 9 | 4.0 | 16.0 | 0.250000 |
| 3 | 39 | 10 | 8 | 1.0 | 18.0 | 0.055556 |
| 4 | 114 | 12 | 21 | 4.0 | 30.0 | 0.133333 |

| | freq_q1+q2 | freq_q1-q2 |
|---|------------|------------|
| 0 | 4 | 0 |
| 1 | 3 | 1 |
| 2 | 35 | 11 |
| 3 | 4 | 0 |
| 4 | 6 | 4 |

[5 rows x 2026 columns]

4.2 Converting strings to numerics

```
[6]: # after we read from sql table each entry was read it as a string
# we convert all the features into numeric before we apply any model
cols = list(X_train.columns)
for i in cols:
    X_train[i] = X_train[i].apply(pd.to_numeric)
    #print(i)
```

```
[7]: # after we read from sql table each entry was read it as a string
# we convert all the features into numeric before we apply any model
cols = list(X_test.columns)
for i in cols:
    X_test[i] = X_test[i].apply(pd.to_numeric)
    #print(i)
```

4.3 Generic method for plotting confusion matrix

```
[8]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i_
    →are predicted class j

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in_
    →that column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to_
    →rows in two dimensional array
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]]
```

```

#                                     [2/3, 4/7]]

# ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
#                               [3/7, 4/7]]
# sum of row elements = 1

B =(C/C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in
→that row
# C = [[1, 2],
#       [3, 4]]
# C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to
→rows in two dimensional array
# C.sum(axis = 0) = [[4, 6]]
# (C/C.sum(axis=0)) = [[1/4, 2/6],
#                       [3/4, 4/6]]
plt.figure(figsize=(20,4))

labels = [1,2]
# representing A in heatmap format
cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
→yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
→yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels,
→yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()


```

4.4 Building a random model (Finding worst-case log-loss)

```
[9]: # we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their
    ↳ sum
# ref: https://stackoverflow.com/a/18662466/4084039
# we create a output array that has exactly same size as the CV data
test_len = y_test.shape[0]
predicted_y = np.zeros((test_len,2))
for i in range(test_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test, predicted_y,
    ↳eps=1e-15))

predicted_y = np.argmax(predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)
```

Log loss on Test Data using Random Model 0.8946451543795837



4.4 Logistic Regression with hyperparameter tuning

```
[10]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/
    ↳ generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15,
    ↳ fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None,
    ↳ learning_rate=optimal, eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with
    ↳ Stochastic Gradient Descent.
# predict(X)          Predict class labels for samples in X.
```

```

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_,
    ↪eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:",log_loss(y_test,
    ↪predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array,c='g')
for i, txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
    ↪random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
    ↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
    ↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y =np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)


```

For values of alpha = 1e-05 The log loss is: 0.42107034776101715
 For values of alpha = 0.0001 The log loss is: 0.4467932754269884

For values of alpha = 0.001 The log loss is: 0.45105223648176035
 For values of alpha = 0.01 The log loss is: 0.4320378625142579
 For values of alpha = 0.1 The log loss is: 0.4492084241860258
 For values of alpha = 1 The log loss is: 0.4894598746560532
 For values of alpha = 10 The log loss is: 0.5432669264172524



For values of best alpha = 1e-05 The train log loss is: 0.41918545363002374
 For values of best alpha = 1e-05 The test log loss is: 0.42107034776101715
 Total number of data points : 60000



4.5 Linear SVM with hyperparameter tuning

```
[11]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.
```

```

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15,
#   →fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None,
#   →learning_rate=optimal, eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with
#   →Stochastic Gradient Descent.
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_,
    →eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test,
    →predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge',
    →random_state=42)

```




```

clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)


predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
→", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
→", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)

```

For values of alpha = 1e-05 The log loss is: 0.4111543493357448
 For values of alpha = 0.0001 The log loss is: 0.424165337237207
 For values of alpha = 0.001 The log loss is: 0.4475273946489084
 For values of alpha = 0.01 The log loss is: 0.459505556906909
 For values of alpha = 0.1 The log loss is: 0.48317805572479033
 For values of alpha = 1 The log loss is: 0.5807274815927826
 For values of alpha = 10 The log loss is: 0.6428145505812121



For values of best alpha = 1e-05 The train log loss is: 0.4072066344844586
 For values of best alpha = 1e-05 The test log loss is: 0.4111543493357448
 Total number of data points : 60000



4.6 XGBoost

4.6.1 Reading data from file

```
[12]: X_train = pd.read_csv('final_features_w2v_train.csv')
      X_test = pd.read_csv('final_features_w2v_test.csv')

[13]: y_train = pd.read_csv('final_features_tfidf_train_label.csv',header=None)
      y_test = pd.read_csv('final_features_tfidf_test_label.csv',header=None)
      y_train = y_train[1]
      y_train = y_train.values.reshape((y_train.shape[0],1))
      print(y_train.shape)
      y_test = y_test[1]
      y_test = y_test.values.reshape((y_test.shape[0],1))
      print(y_test.shape)
```

(140000, 1)

(60000, 1)

```
[14]: X_train.drop(['Unnamed: 0', 'id','is_duplicate'], axis=1, inplace=True)
      X_test.drop(['Unnamed: 0', 'id','is_duplicate'], axis=1, inplace=True)
```

```
[15]: X_train.head()
```

```
[15]:
```

| | 0_x | 1_x | 2_x | 3_x | 4_x | 5_x | \ |
|---|------------|-----------|------------|------------|------------|------------|---|
| 0 | 2.439142 | 18.114910 | -12.255746 | 104.130825 | -19.363592 | -2.205739 | |
| 1 | -15.773163 | 36.496838 | 8.192818 | 47.599665 | 37.109427 | 62.937893 | |
| 2 | -12.001888 | 88.561495 | 54.594788 | 1.760126 | -36.467282 | 33.917674 | |
| 3 | -18.628788 | 74.915513 | 115.482397 | 50.463298 | -12.816265 | -43.315368 | |
| 4 | -2.920037 | 64.090752 | 9.972080 | 105.365901 | 35.781809 | 31.707977 | |

| | 6_x | 7_x | 8_x | 9_x | ... | freq_qid2 | q1len | q2len | \ |
|---|-------------|-------------|-----------|------------|-----|-----------|-------|-------|---|
| 0 | -13.635292 | -3.182549 | 30.811532 | -4.650403 | ... | 2 | 40 | 33 | |
| 1 | -38.917103 | -61.323661 | 5.005070 | -25.191947 | ... | 1 | 50 | 35 | |
| 2 | -59.055783 | -20.182544 | 43.250478 | -36.022031 | ... | 12 | 37 | 48 | |
| 3 | -102.989147 | -104.574963 | 46.028487 | -52.984333 | ... | 2 | 60 | 39 | |
| 4 | -89.826797 | -13.955268 | 27.267864 | 18.659436 | ... | 1 | 59 | 114 | |

| | q1_n_words | q2_n_words | word_Common | word_Total | word_share | freq_q1+q2 | \ |
|---|------------|------------|-------------|------------|------------|------------|---|
| 0 | 9 | 7 | 6.0 | 16.0 | 0.375000 | 4 | |

| | | | | | | |
|---|----|----|-----|------|----------|----|
| 1 | 9 | 8 | 6.0 | 16.0 | 0.375000 | 3 |
| 2 | 7 | 9 | 4.0 | 16.0 | 0.250000 | 35 |
| 3 | 10 | 8 | 1.0 | 18.0 | 0.055556 | 4 |
| 4 | 12 | 21 | 4.0 | 30.0 | 0.133333 | 6 |

| | freq_q1-q2 |
|---|------------|
| 0 | 0 |
| 1 | 1 |
| 2 | 11 |
| 3 | 0 |
| 4 | 4 |

[5 rows x 794 columns]

4.6.2 Converting strings to numerics

```
[16]: # after we read from sql table each entry was read it as a string
# we convert all the features into numeric before we apply any model
cols = list(X_train.columns)
for i in cols:
    X_train[i] = X_train[i].apply(pd.to_numeric)
    #print(i)
```

```
[17]: # after we read from sql table each entry was read it as a string
# we convert all the features into numeric before we apply any model
cols = list(X_test.columns)
for i in cols:
    X_test[i] = X_test[i].apply(pd.to_numeric)
    #print(i)
```

```
[18]: from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBClassifier
params = {
    'min_child_weight': [1, 5, 10],
    'gamma': [0.5, 1, 1.5, 2, 5],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'max_depth': [3, 4, 5],
    'eta': [0.02, 0.01, 0.1]
}

xgb = XGBClassifier(learning_rate=0.02, n_estimators=400, objective='binary:
    ↳logistic',
                    silent=True, early_stopping_rounds=20)
```

```
[19]: folds = 3
param_comb = 10
skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)
```

```

random_search = RandomizedSearchCV(xgb, param_distributions=params,
    ↪n_iter=param_comb, scoring='neg_log_loss', n_jobs=4, cv=skf.split(X_train,
    ↪y_train), verbose=3, random_state=1001 )
history = random_search.fit(X_train, y_train)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 30 out of 30 | elapsed: 220.1min finished

```

```
[20]: random_search.best_params_
```

```

[20]: {'subsample': 0.6,
      'min_child_weight': 1,
      'max_depth': 5,
      'gamma': 2,
      'eta': 0.02,
      'colsample_bytree': 1.0}

```

```

[21]: import xgboost as xgb
      params = {}
      params['objective'] = 'binary:logistic'
      params['eval_metric'] = 'logloss'
      params['eta'] = 0.02
      params['max_depth'] = 5
      params['booster'] = 'gbtree'
      params['gama'] = 2
      params['max_depth'] = 5
      params['min_child_weight'] = 1
      params['subsample'] = 0.6
      params['colsample_bytree'] = 1

      d_train = xgb.DMatrix(X_train, label=y_train)
      d_test = xgb.DMatrix(X_test, label=y_test)

      watchlist = [(d_train, 'train'), (d_test, 'valid')]

      bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20,
    ↪verbose_eval=10)

      xgdmatrix = xgb.DMatrix(X_train, y_train)
      predict_y = bst.predict(d_test)
      print("The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_,
    ↪eps=1e-15))

```

```
[0]    train-logloss:0.683686  valid-logloss:0.684189
```

Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.


Will train until valid-logloss hasn't improved in 20 rounds.

| | | |
|-------|------------------------|------------------------|
| [10] | train-logloss:0.610007 | valid-logloss:0.610489 |
| [20] | train-logloss:0.556155 | valid-logloss:0.557228 |
| [30] | train-logloss:0.515604 | valid-logloss:0.516966 |
| [40] | train-logloss:0.484679 | valid-logloss:0.486367 |
| [50] | train-logloss:0.460706 | valid-logloss:0.462706 |
| [60] | train-logloss:0.44181 | valid-logloss:0.44415 |
| [70] | train-logloss:0.426772 | valid-logloss:0.429376 |
| [80] | train-logloss:0.41461 | valid-logloss:0.417434 |
| [90] | train-logloss:0.404625 | valid-logloss:0.407675 |
| [100] | train-logloss:0.396265 | valid-logloss:0.399621 |
| [110] | train-logloss:0.389444 | valid-logloss:0.393158 |
| [120] | train-logloss:0.383771 | valid-logloss:0.387708 |
| [130] | train-logloss:0.378983 | valid-logloss:0.383192 |
| [140] | train-logloss:0.374681 | valid-logloss:0.379133 |
| [150] | train-logloss:0.370981 | valid-logloss:0.375716 |
| [160] | train-logloss:0.367785 | valid-logloss:0.372786 |
| [170] | train-logloss:0.364977 | valid-logloss:0.370218 |
| [180] | train-logloss:0.362547 | valid-logloss:0.368006 |
| [190] | train-logloss:0.360328 | valid-logloss:0.366092 |
| [200] | train-logloss:0.358146 | valid-logloss:0.364216 |
| [210] | train-logloss:0.356406 | valid-logloss:0.36272 |
| [220] | train-logloss:0.354673 | valid-logloss:0.36121 |
| [230] | train-logloss:0.353141 | valid-logloss:0.360046 |
| [240] | train-logloss:0.351529 | valid-logloss:0.358695 |
| [250] | train-logloss:0.349967 | valid-logloss:0.357449 |
| [260] | train-logloss:0.348429 | valid-logloss:0.356223 |
| [270] | train-logloss:0.346931 | valid-logloss:0.35514 |
| [280] | train-logloss:0.345593 | valid-logloss:0.354215 |
| [290] | train-logloss:0.344015 | valid-logloss:0.352978 |
| [300] | train-logloss:0.342692 | valid-logloss:0.351988 |
| [310] | train-logloss:0.341379 | valid-logloss:0.351134 |
| [320] | train-logloss:0.340107 | valid-logloss:0.350297 |
| [330] | train-logloss:0.338866 | valid-logloss:0.349517 |
| [340] | train-logloss:0.337701 | valid-logloss:0.348788 |
| [350] | train-logloss:0.336542 | valid-logloss:0.348085 |
| [360] | train-logloss:0.335393 | valid-logloss:0.347365 |
| [370] | train-logloss:0.33424 | valid-logloss:0.346627 |
| [380] | train-logloss:0.333213 | valid-logloss:0.346047 |
| [390] | train-logloss:0.332199 | valid-logloss:0.345358 |
| [399] | train-logloss:0.33119 | valid-logloss:0.34472 |

The test log loss is: 0.34471666093470993

```
[22]: predicted_y = np.array(predict_y>0.5,dtype=int)
      print("Total number of data points :", len(predicted_y))
      plot_confusion_matrix(y_test, predicted_y)
```

Total number of data points : 60000



```
[24]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "train log loss", "test log loss", "train vs test precision", "train vs test recall"]
x.add_row(["TFIDF", "Linear Regression", 0.4191, 0.429, "0.813 - 0.760", "0.878 - 0.656"])
x.add_row(["TFIDF", "SVM", 0.4072, 0.4111, "0.817 - 0.769", "0.883 - 0.664"])
x.add_row(["TFIDFAverageW2V", "XGBoost", 0.3311, 0.3447, "0.843 - 0.802", "0.896 - 0.716"])
x.border=True
print(x)
```

| Vectorizer | Model | train log loss | test log loss | train vs test precision | train vs test recall |
|-----------------|-------------------|----------------|---------------|-------------------------|----------------------|
| TFIDF | Linear Regression | 0.4191 | 0.429 | 0.813 - 0.760 | 0.878 - 0.656 |
| TFIDF | SVM | 0.4072 | 0.4111 | 0.817 - 0.769 | 0.883 - 0.664 |
| TFIDFAverageW2V | XGBoost | 0.3311 | 0.3447 | 0.843 - 0.802 | 0.896 - 0.716 |

Clearly XGBoost Model is the best model as it has minimum train and test log loss and precision and recall are better than other two models